

CS224N Assignment 1

Zheng Nie

TOTAL POINTS

24 / 25

QUESTION 1

Question 1 10 pts

1.1 distinct_words [coding] 2 / 2

✓ + 2 pts Correct

+ 0 pts Incorrect

1.2 compute_co_occurrence_matrix [coding] 3 / 3

✓ - 0 pts Correct

- 3 pts Incorrect or missing

1.3 reduce_to_k_dim [coding] 1 / 1

✓ + 1 pts Passed all tests

+ 0 pts Code incorrect

1.4 plot_embeddings [coding] 1 / 1

✓ + 1 pts Correct

+ 0 pts Incorrect/Unattempted

1.5 co-occurrence plot analysis [written] 3 / 3

✓ + 1 pts Correct plot

✓ + 1 pts Correctly noting what clusters together

✓ + 1 pts Correctly noting what does not cluster together

+ 0 pts Missing/incorrect plot

QUESTION 2

Question 2 15 pts

2.1 word2vec plot analysis 4 / 4

✓ + 1 pts Correct Plot

✓ + 1 pts Clusters pointed out correctly

✓ + 1 pts Non-cluster points pointed out correctly

✓ + 1 pts Pointed out what's different

+ 0 pts Incorrect Plot/Unattempted

2.2 polysemous words 2 / 2

✓ + 1 pts Valid polysemous word with multiple meanings in top 10

✓ + 1 pts Plausible possible reason for why the

chosen word's meanings show up in top 10 or why other polysemous words didn't work.

+ 0 pts No answer.

2.3 synonyms & antonyms 2 / 2

✓ - 0 pts Correct

- 1 pts Did not provide a possible reason

- 1 pts Did not provide valid triple

2.4 correct analogy 2 / 2

✓ + 2 pts Valid analogy is top-1

+ 1 pts Valid analogy but not top-1

+ 1 pts Unclear analogy

+ 0 pts Incorrect/missing analogy

2.5 incorrect analogy 0 / 1

+ 1 pts Correct analogy, top-1 incorrect

✓ + 0 pts Incorrect analogy

+ 0 pts Correct analogy, top-1 correct

+ 0 pts Correct analogy, incorrect

implementation/code

+ 0 pts Not attempted

2.6 guided bias exploration 1 / 1

✓ - 0 pts Correct

- 1 pts Incorrect / Insufficient explanation

2.7 independent bias exploration 2 / 2

✓ + 1 pts Found bias

✓ + 1 pts Explained bias

+ 0 pts Incorrect/Not attempted

+ 1 pts 2-way comparison not provided and hence cannot comment on bias.

2.8 thinking about bias 1 / 1

✓ + 1 pts Correct: biases are picked up from the training corpus (Google News)

+ 0 pts Insufficient/missing explanation

```
In [8]: def distinct_words(corpus):
    """ Determine a list of distinct words for the corpus.
    Params:
        corpus (list of list of strings): corpus of documents
    Return:
        corpus_words (list of strings): list of distinct words across the corpus, sorted (using python 'sorted' function)
        num_corpus_words (integer): number of distinct words across the corpus
    """
    corpus_words = []
    num_corpus_words = -1

    # -----
    # Write your implementation here.
    corpus_words = sorted(list(set(j for i in corpus for j in i)))
    num_corpus_words = len(corpus_words)

    # -----
    return corpus_words, num_corpus_words
```

```
In [9]: # -----
# Run this sanity check
# Note that this not an exhaustive check for correctness.
# -----

# Define toy corpus
test_corpus = ["START All that glitters isn't gold END".split(" "),
               "START All's well that ends well END".split(" ")]
test_corpus_words, num_corpus_words = distinct_words(test_corpus)

# Correct answers
ans_test_corpus_words = sorted(list(set(["START", "All", "ends", "tha",
                                         "t", "gold", "All's", "glitters", "isn't", "well", "END"])))
ans_num_corpus_words = len(ans_test_corpus_words)

# Test correct number of words
assert (num_corpus_words == ans_num_corpus_words), "Incorrect number of distinct words. Correct: {}. Yours: {}".format(ans_num_corpus_words, num_corpus_words)

# Test correct words
assert (test_corpus_words == ans_test_corpus_words), "Incorrect corpus words.\nCorrect: {}\nYours: {}".format(str(ans_test_corpus_words), str(test_corpus_words))

# Print Success
print("-" * 80)
print("Passed All Tests!")
print("-" * 80)
```


Passed All Tests!

Question 1.2: Implement compute_co_occurrence_matrix [code] (3 points)

Write a method that constructs a co-occurrence matrix for a certain window-size n (with a default of 4), considering words n before and n after the word in the center of the window. Here, we start to use `numpy` (`np`) to represent vectors, matrices, and tensors. If you're not familiar with NumPy, there's a NumPy tutorial in the second half of this cs231n [Python NumPy tutorial](http://cs231n.github.io/python-numpy-tutorial/) (<http://cs231n.github.io/python-numpy-tutorial/>).

1.1 distinct_words [coding] 2 / 2

✓ + 2 pts Correct

+ 0 pts Incorrect

```
In [10]: def compute_co_occurrence_matrix(corpus, window_size=4):
    """ Compute co-occurrence matrix for the given corpus and window_
size (default of 4).

    Note: Each word in a document should be at the center of a wi
ndow. Words near edges will have a smaller
    number of co-occurring words.

    For example, if we take the document "START All that gl
itters is not gold END" with window size of 4,
        "All" will co-occur with "START", "that", "glitters",
    "is", and "not".

    Params:
        corpus (list of list of strings): corpus of documents
        window_size (int): size of context window
    Return:
        M (numpy matrix of shape (number of corpus words, number
of number of corpus words)):
            Co-occurrence matrix of word counts.
            The ordering of the words in the rows/columns should
be the same as the ordering of the words given by the distinct_words
function.
        word2Ind (dict): dictionary that maps word to index (i.e.
row/column number) for matrix M.
    """
    words, num_words = distinct_words(corpus)
    M = None

    word2Ind = {}

    # -----
    # Write your implementation here.
    M = np.zeros((num_words, num_words))

    # generate word2Ind for M indexing
    for i, word in enumerate(words):
        word2Ind[word] = i

    for corp in corpus:
        for s, sub_string in enumerate(corp):
            sub_string_word_to_check = corp[0 if s - window_size < 0 el
se s - window_size : s]
            sub_string_word_to_check.extend(corp[s + 1: s + window_size + 1])
            i = word2Ind[sub_string]
            for string in sub_string_word_to_check:
                M[i][word2Ind[string]] = M[i][word2Ind[string]] + 1.0

    #
    return M, word2Ind
```

```
In [11]: # -----
# Run this sanity check
# Note that this is not an exhaustive check for correctness.
# -----
```

```
# Define toy corpus and get student's co-occurrence matrix
test_corpus = ["START All that glitters isn't gold END".split(" "),
               "START All's well that ends well END".split(" ")]
M_test, word2Ind_test = compute_co_occurrence_matrix(test_corpus, window_size=1)

# Correct M and word2Ind
M_test_ans = np.array(
    [[0., 0., 0., 1., 0., 0., 0., 1., 0.,],
     [0., 0., 0., 1., 0., 0., 0., 0., 1.,],
     [0., 0., 0., 0., 0., 1., 0., 0., 1.,],
     [1., 1., 0., 0., 0., 0., 0., 0., 0.,],
     [0., 0., 0., 0., 0., 0., 0., 1., 1.,],
     [0., 0., 0., 0., 0., 0., 1., 1., 0.,],
     [0., 0., 1., 0., 0., 0., 1., 0., 0.,],
     [0., 0., 0., 0., 1., 1., 0., 0., 0.,],
     [1., 0., 0., 0., 1., 1., 0., 0., 1.,],
     [0., 1., 1., 0., 1., 0., 0., 1., 0.,]])
)
word2Ind_ans = {'All': 0, 'All's': 1, 'END': 2, 'START': 3, 'ends': 4,
                 'glitters': 5, 'gold': 6, "isn't": 7, 'that': 8, 'well': 9}

# Test correct word2Ind
assert (word2Ind_ans == word2Ind_test), "Your word2Ind is incorrect:\n\nCorrect: {}\nYours: {}".format(word2Ind_ans, word2Ind_test)

# Test correct M shape
assert (M_test.shape == M_test_ans.shape), "M matrix has incorrect shape.\n\nCorrect: {}\nYours: {}".format(M_test.shape, M_test_ans.shape)

# Test correct M values
for w1 in word2Ind_ans.keys():
    idx1 = word2Ind_ans[w1]
    for w2 in word2Ind_ans.keys():
        idx2 = word2Ind_ans[w2]
        student = M_test[idx1, idx2]
        correct = M_test_ans[idx1, idx2]
        if student != correct:
            print("Correct M:")
            print(M_test_ans)
            print("Your M: ")
            print(M_test)
            raise AssertionError("Incorrect count at index ({}, {})=( {}, {}) in matrix M. Yours has {} but should have {}.".format(idx1, idx2, w1, w2, student, correct))

# Print Success
print ("-" * 80)
print("Passed All Tests!")
print ("-" * 80)
```

Passed All Tests!

Question 1.3: Implement `reduce_to_k_dim` [code] (1 point)

Construct a method that performs dimensionality reduction on the matrix to produce k-dimensional embeddings. Use SVD to take the top k components and produce a new matrix of k-dimensional embeddings.

Note: All of numpy, scipy, and scikit-learn (`sklearn`) provide some implementation of SVD, but only scipy and sklearn provide an implementation of Truncated SVD, and only sklearn provides an efficient randomized algorithm for calculating large-scale Truncated SVD. So please use [`sklearn.decomposition.TruncatedSVD`](https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html) (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html>).

```
In [12]: def reduce_to_k_dim(M, k=2):
    """ Reduce a co-occurrence count matrix of dimensionality (num_corpus_words, num_corpus_words)
        to a matrix of dimensionality (num_corpus_words, k) using the
        following SVD function from Scikit-Learn:
            - http://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html

    Params:
        M (numpy matrix of shape (number of corpus words, number
        of number of corpus words)): co-occurrence matrix of word counts
        k (int): embedding size of each word after dimension reduction

    Return:
        M_reduced (numpy matrix of shape (number of corpus words,
        k)): matrix of k-dimensionaoal word embeddings.
                In terms of the SVD from math class, this actuall
        y returns U * S
    """
    n_iters = 10      # Use this parameter in your call to `TruncatedS
    VD`  

    M_reduced = None  

    print("Running Truncated SVD over %i words..." % (M.shape[0]))  

    # -----  

    # Write your implementation here.  

    svd = TruncatedSVD(n_components=k, n_iter=n_iters)  

    M_reduced = svd.fit_transform(M)  

    # -----  

    print("Done.")  

    return M_reduced
```

1.2 compute_co_occurrence_matrix [coding] 3 / 3

✓ - 0 pts Correct

- 3 pts Incorrect or missing

Passed All Tests!

Question 1.3: Implement `reduce_to_k_dim` [code] (1 point)

Construct a method that performs dimensionality reduction on the matrix to produce k-dimensional embeddings. Use SVD to take the top k components and produce a new matrix of k-dimensional embeddings.

Note: All of numpy, scipy, and scikit-learn (`sklearn`) provide some implementation of SVD, but only scipy and sklearn provide an implementation of Truncated SVD, and only sklearn provides an efficient randomized algorithm for calculating large-scale Truncated SVD. So please use [`sklearn.decomposition.TruncatedSVD`](https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html) (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html>).

```
In [12]: def reduce_to_k_dim(M, k=2):
    """ Reduce a co-occurrence count matrix of dimensionality (num_corpus_words, num_corpus_words)
        to a matrix of dimensionality (num_corpus_words, k) using the
        following SVD function from Scikit-Learn:
            - http://scikit-learn.org/stable/modules/generated/sklearn.decomposition.TruncatedSVD.html

    Params:
        M (numpy matrix of shape (number of corpus words, number
        of number of corpus words)): co-occurrence matrix of word counts
        k (int): embedding size of each word after dimension reduction

    Return:
        M_reduced (numpy matrix of shape (number of corpus words,
        k)): matrix of k-dimensionaoal word embeddings.
            In terms of the SVD from math class, this actuall
        y returns U * S
    """
    n_iters = 10      # Use this parameter in your call to `TruncatedS
    VD`  

    M_reduced = None  

    print("Running Truncated SVD over %i words..." % (M.shape[0]))  

    # -----  

    # Write your implementation here.  

    svd = TruncatedSVD(n_components=k, n_iter=n_iters)  

    M_reduced = svd.fit_transform(M)  

    # -----  

    print("Done.")  

    return M_reduced
```

```
In [13]: # -----
# Run this sanity check
# Note that this not an exhaustive check for correctness
# In fact we only check that your M_reduced has the right dimensions.
# ----- 

# Define toy corpus and run student code
test_corpus = ["START All that glitters isn't gold END".split(" "),
                "START All's well that ends well END".split(" ")]
M_test, word2Ind_test = compute_co_occurrence_matrix(test_corpus, window_size=1)
M_test_reduced = reduce_to_k_dim(M_test, k=2)

# Test proper dimensions
assert (M_test_reduced.shape[0] == 10), "M_reduced has {} rows; should have {}".format(M_test_reduced.shape[0], 10)
assert (M_test_reduced.shape[1] == 2), "M_reduced has {} columns; should have {}".format(M_test_reduced.shape[1], 2)

# Print Success
print ("-" * 80)
print("Passed All Tests!")
print ("-" * 80)
```

Running Truncated SVD over 10 words...
Done.

Passed All Tests!

Question 1.4: Implement `plot_embeddings` [code] (1 point)

Here you will write a function to plot a set of 2D vectors in 2D space. For graphs, we will use Matplotlib (`plt`).

For this example, you may find it useful to adapt [this code](#)

(<https://www.pythontutorial.net/2018/05/08/matplotlib-scatter-plot-annotate-set-text-at-label-each-point/>). In the future, a good way to make a plot is to look at [the Matplotlib gallery](https://matplotlib.org/gallery/index.html) (<https://matplotlib.org/gallery/index.html>), find a plot that looks somewhat like what you want, and adapt the code they give.

1.3 reduce_to_k_dim [coding] 1 / 1

✓ + 1 pts Passed all tests

+ 0 pts Code incorrect

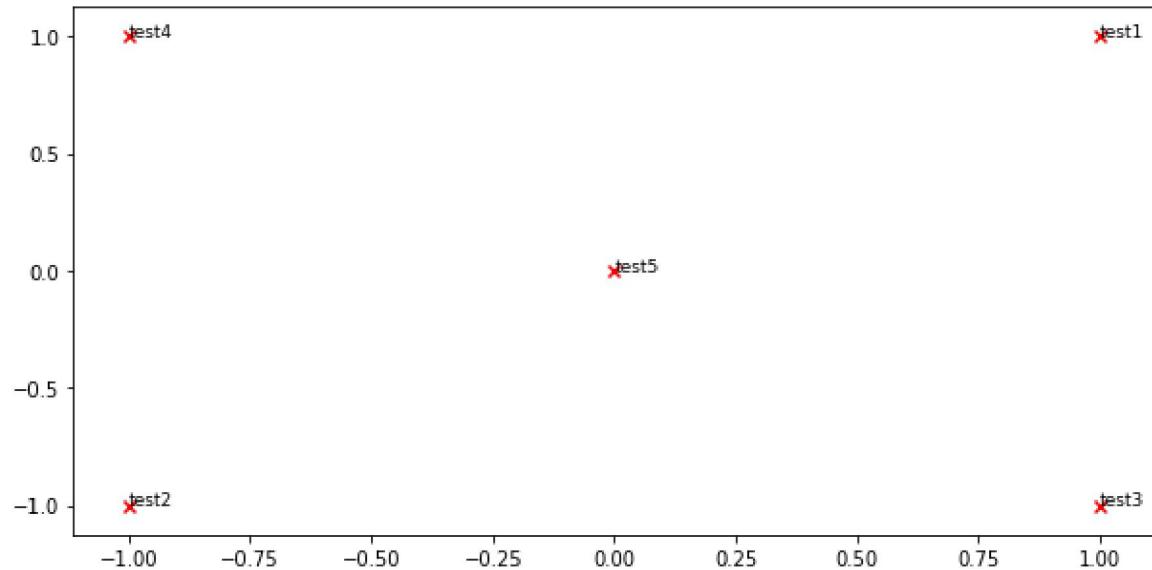
```
In [14]: def plot_embeddings(M_reduced, word2Ind, words):
    """ Plot in a scatterplot the embeddings of the words specified in the list "words".
        NOTE: do not plot all the words listed in M_reduced / word2Ind.
        Include a label next to each point.

    Params:
        M_reduced (numpy matrix of shape (number of corpus words, k)): matrix of k-dimensioal word embeddings
        word2Ind (dict): dictionary that maps word to indices for matrix M
        words (list of strings): words whose embeddings we want to visualize
    """
    # -----
    # Write your implementation here.
    for word in words:
        coords = M_reduced[word2Ind[word]]
        x = coords[0]
        y = coords[1]
        plt.scatter(x, y, marker='x', color='red')
        plt.text(x, y, word, fontsize=9)
    plt.show()

    # -----
```

```
In [15]: # -----
# Run this sanity check
# Note that this not an exhaustive check for correctness.
# The plot produced should look like the "test solution plot" depicted below.
# -----  
  
print ("-" * 80)
print ("Outputted Plot:")  
  
M_reduced_plot_test = np.array([[1, 1], [-1, -1], [1, -1], [-1, 1], [0, 0]])
word2Ind_plot_test = {'test1': 0, 'test2': 1, 'test3': 2, 'test4': 3,
                      'test5': 4}
words = ['test1', 'test2', 'test3', 'test4', 'test5']
plot_embeddings(M_reduced_plot_test, word2Ind_plot_test, words)  
  
print ("-" * 80)
```


Outputted Plot:



1.4 plot_embeddings [coding] 1 / 1

✓ + 1 pts Correct

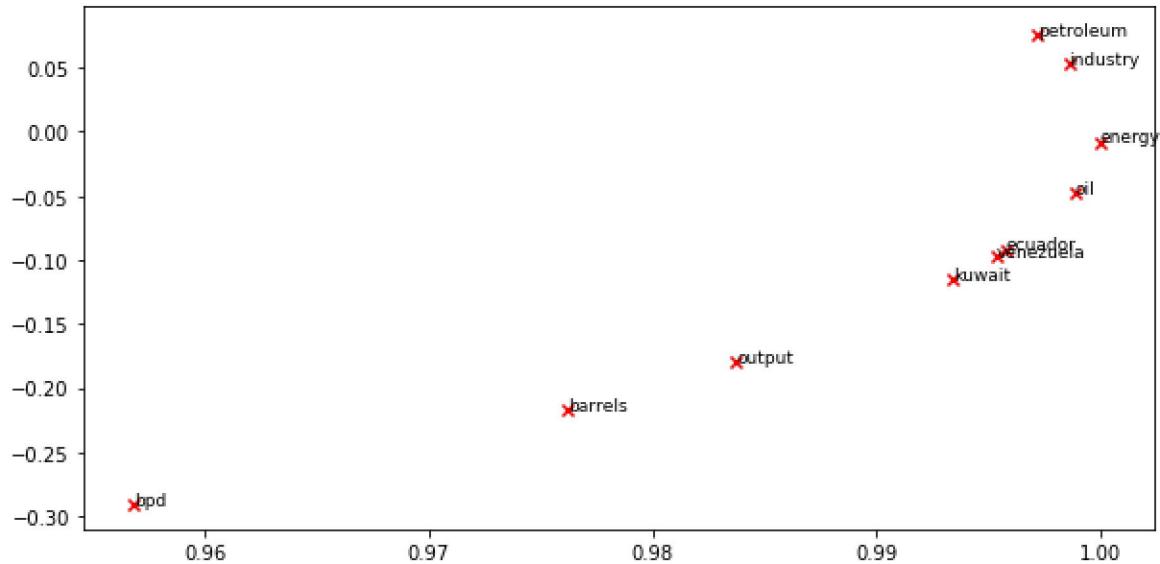
+ 0 pts Incorrect/Unattempted

```
In [16]: # -----
# Run This Cell to Produce Your Plot
# -----
reuters_corpus = read_corpus()
M_co_occurrence, word2Ind_co_occurrence = compute_co_occurrence_matrix(reuters_corpus)
M_reduced_co_occurrence = reduce_to_k_dim(M_co_occurrence, k=2)

# Rescale (normalize) the rows to make them each of unit-length
M_lengths = np.linalg.norm(M_reduced_co_occurrence, axis=1)
M_normalized = M_reduced_co_occurrence / M_lengths[:, np.newaxis] # broadcasting

words = ['barrels', 'bpd', 'ecuador', 'energy', 'industry', 'kuwait',
         'oil', 'output', 'petroleum', 'venezuela']
plot_embeddings(M_normalized, word2Ind_co_occurrence, words)
```

Running Truncated SVD over 8185 words...
Done.



Answer:

kuwait,venezuela,ecuador which are all countries (names) have clustered together. Those countries also close to oil, since those country produces a lot of oil. energy and oil are closed since oil can generate genergy so it is related to energy.Petroleum and industry are closed since petroleum is kind of industry,so industry is generalizion of petroleum.

oil and petroleum should clustered, since they have similar meaning. bpd and output should have clustered as they could have samilar meaning of the productitity outputs of petroleum industry

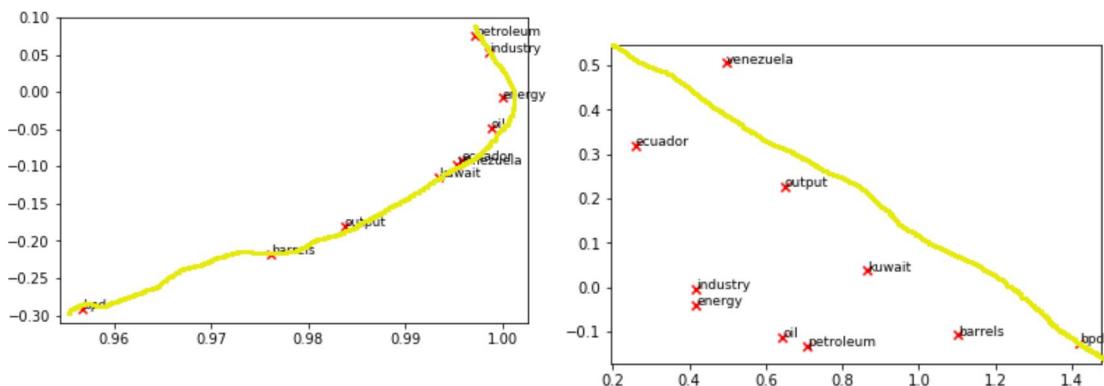
1.5 co-occurrence plot analysis [written] 3 / 3

- ✓ + 1 pts Correct plot
- ✓ + 1 pts Correctly noting what clusters together
- ✓ + 1 pts Correctly noting what does not cluster together
- + 0 pts Missing/incorrect plot

Answer:

oil and petroleum which have similar meaning are clustered. industry and energy are clustered. ecuador, kuwait, venezuela which are all petroleum production country name should have clustered, output and bpd have similar meaning for petroleum production amount should have clustered

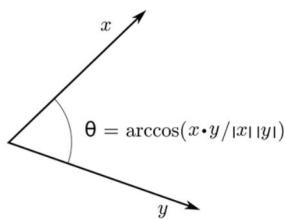
The difference this plot from the one generated earlier is: First, the scale of X, Y axes is different. Second, the distribution in plot for same word is different between those 2 plot, in the plot generate earlier we could see all the words distribution is along/closed to a curved line, for the second one we just generated is spread in the left down side of the plot(See following two figures). Third different is the evaluation result is different. in the plot generated earlier shows those oil countries names are clustered together and close to oil. in the new one, we didn't see this. Forth, for same word in those 2 word embedding plot the word vector weight/value (coordinates) is different. For example, for oil in first plot, it is (1.00, -0.05), the oil in second plot has (0.62, -0.11)



Cosine Similarity

Now that we have word vectors, we need a way to quantify the similarity between individual words, according to these vectors. One such metric is cosine-similarity. We will be using this to find words that are "close" and "far" from one another.

We can think of n-dimensional vectors as points in n-dimensional space. If we take this perspective L1 and L2 Distances help quantify the amount of space "we must travel" to get between these two points. Another approach is to examine the angle between two vectors. From trigonometry we know that:



Instead of computing the actual angle, we can leave the similarity in terms of $similarity = \cos(\Theta)$. Formally the [Cosine Similarity \(\[https://en.wikipedia.org/wiki/Cosine_similarity\]\(https://en.wikipedia.org/wiki/Cosine_similarity\)\)](https://en.wikipedia.org/wiki/Cosine_similarity), s between two vectors p and q is defined as:

$$s = \frac{\mathbf{p} \cdot \mathbf{q}}{\|\mathbf{p}\| \|\mathbf{q}\|}, \text{ where } s \in [-1, 1]$$

2.1 word2vec plot analysis 4 / 4

- ✓ + 1 pts Correct Plot
- ✓ + 1 pts Clusters pointed out correctly
- ✓ + 1 pts Non-cluster points pointed out correctly
- ✓ + 1 pts Pointed out what's different
- + 0 pts Incorrect Plot/Unattempted

Question 2.2: Polysemous Words (2 points) [code + written]

Find a [polysemous](https://en.wikipedia.org/wiki/Polysemy) (<https://en.wikipedia.org/wiki/Polysemy>) word (for example, "leaves" or "scoop") such that the top-10 most similar words (according to cosine similarity) contains related words from *both* meanings. For example, "leaves" has both "vanishes" and "stalks" in the top 10, and "scoop" has both "handed_waffle_cone" and "lowdown". You will probably need to try several polysemous words before you find one. Please state the polysemous word you discover and the multiple meanings that occur in the top 10. Why do you think many of the polysemous words you tried didn't work?

Note: You should use the `wv_from_bin.most_similar(word)` function to get the top 10 similar words. This function ranks all other words in the vocabulary with respect to their cosine similarity to the given word. For further assistance please check the [GenSim documentation](https://radimrehurek.com/gensim/models/keyedvectors.html#gensim.models.KeyedVectors.FastTextKeyer) (<https://radimrehurek.com/gensim/models/keyedvectors.html#gensim.models.KeyedVectors.FastTextKeyer>)

```
In [23]: # -----
# Write your polysemous word exploration code here.
```

```
wv_from_bin.most_similar("left")
# -----
```

```
Out[23]: [('leaving', 0.6707000732421875),
('leave', 0.525093138217926),
('leaves', 0.5228644609451294),
('returned', 0.5059226751327515),
('right', 0.49213993549346924),
('departed', 0.49109700322151184),
('limped', 0.48599502444267273),
('went', 0.4719873070716858),
('remaining', 0.4650370478630066),
('empty', 0.4546155333518982)]
```

Answer:

The word I choose is "left", which has "leave" and "remaining".

Most of the meanings of the polysemous have are unrelated, it is possible that the meaning cluster for those meanings in the embedding spaces are not close to each other. and the number of similar words for each cluster could have more than the polysemous words and more closer to each other in word embedding spaces. which also means the one or some polysemous in the corpus traing the wv are more often used or "contexted" than others. and also only pick top 10, so it even reduces the chance to show the simility for polysemous in the list.

2.2 polysemous words 2 / 2

- ✓ + 1 pts Valid polysemous word with multiple meanings in top 10
- ✓ + 1 pts Plausible possible reason for why the chosen word's meanings show up in top 10 or why other polysemous words didn't work.
- + 0 pts No answer.

Question 2.3: Synonyms & Antonyms (2 points) [code + written]

When considering Cosine Similarity, it's often more convenient to think of Cosine Distance, which is simply $1 - \text{Cosine Similarity}$.

Find three words (w_1, w_2, w_3) where w_1 and w_2 are synonyms and w_1 and w_3 are antonyms, but $\text{Cosine Distance}(w_1, w_3) < \text{Cosine Distance}(w_1, w_2)$. For example, $w_1 = \text{"happy"}$ is closer to $w_3 = \text{"sad"}$ than to $w_2 = \text{"cheerful"}$.

Once you have found your example, please give a possible explanation for why this counter-intuitive result may have happened.

You should use the the `wv_from_bin.distance(w1, w2)` function here in order to compute the cosine distance between two words. Please see the [GenSim documentation](https://radimrehurek.com/gensim/models/keyedvectors.html#gensim.models.keyedvectors.FastTextKeyer) (<https://radimrehurek.com/gensim/models/keyedvectors.html#gensim.models.keyedvectors.FastTextKeyer>) for further assistance.

```
In [24]: # -----
# Write your synonym & antonym exploration code here.

w1 = "large"
w2 = "gaint"
w3 = "small"
w1_w2_dist = wv_from_bin.distance(w1, w2)
w1_w3_dist = wv_from_bin.distance(w1, w3)

print("Synonyms {}, {} have cosine distance: {}".format(w1, w2, w1_w2_dist))
print("Antonyms {}, {} have cosine distance: {}".format(w1, w3, w1_w3_dist))

# -----
```

Synonyms large, gaint have cosine distance: 0.8168608269597246
Antonyms large, small have cosine distance: 0.26688501049911195

Answer

The word I found is large, gaint, small.

Here's the reason why this counter-intuitive result may have happened:
 if we run `wv_from_bin.most_similar("large")` like question 2.2, we will get:

```
[('sizeable', 0.7341436147689819),
 ('small', 0.733115017414093),
 ('sizable', 0.7325241565704346),
 ('Large', 0.6654301285743713),
 ('huge', 0.6589166522026062),
 ('larger', 0.6517370939254761),
 ('massive', 0.6127927899360657),
 ('smaller', 0.6025481224060059),
 ('substantial', 0.5873902440071106),
 ('gigantic', 0.5858762860298157)] ,
```

if we run `wv_from_bin.most_similar("small")` we will get

```
[('large', 0.733115017414093),
 ('tiny', 0.718792736530304),
 ('medium_sized', 0.6426598429679871),
 ('Small', 0.6270469427108765),
 ('smaller', 0.619428813457489),
 ('minuscule', 0.5676835775375366),
 ('larger', 0.5446805953979492),
 ('mid_sized', 0.5328439474105835),
 ('midsized', 0.519831120967865),
 ('sizable', 0.5175076723098755)],
```

if we run `wv_from_bin.most_similar("gaint")`, we will get

```
[('giant', 0.5869466066360474),
 ('BOMBAY_XFN_ASIA', 0.4832102656364441),
 ('jt_venture', 0.4771811068058014),
 ('Videocon_Industries_Ltd.', 0.4737812578678131)
 ('TPV_Technology_Ltd', 0.46327781677246094),
 ('outsourcer_Wipro', 0.45941978693008423),
 ('RELI.BO', 0.45818576216697693),
 ('Ltd._India', 0.4561852216720581),
 ('Jindal_Saw_Ltd', 0.4555988311767578),
 ('retailer_Gome', 0.45485517382621765)]
```

For result above , we could see for large and small, the top 10 words are almost same and most focusing on the size property description. for gaint, which is also size property desction word, expecting showing similar result, but top 10 is showing most of them are company names. which means in the word embedding space, word

vector for word "gaint" are closer to the company name clusters. and word vector for "large" and "small" are closer to size descriptions clusters. which also means in the corpus for traing w2v , the gaint are most often being used for comanies or it self represents big companines rather than being used for size.

Solving Analogies with Word Vectors

Word2Vec vectors have been shown to *sometimes* exhibit the ability to solve analogies.

As an example, for the analogy "man : king :: woman : x", what is x?

In the cell below, we show you how to use word vectors to find x. The `most_similar` function finds words that are most similar to the words in the `positive` list and most dissimilar from the words in the `negative` list. The answer to the analogy will be the word ranked most similar (largest numerical value).

Note: Further Documentation on the `most_similar` function can be found within the [GenSim documentation](#) (<https://radimrehurek.com/gensim/models/keyedvectors.html#gensim.models.KeyedVectors.FastTextKeyer>)

```
In [25]: # Run this cell to answer the analogy -- man : king :: woman : x
pprint.pprint(wv_from_bin.most_similar(positive=['woman', 'king'], negative=['man']))
```

```
[('queen', 0.7118192315101624),
 ('monarch', 0.6189675331115723),
 ('princess', 0.5902431011199951),
 ('crown_prince', 0.5499460697174072),
 ('prince', 0.5377321243286133),
 ('kings', 0.5236844420433044),
 ('Queen_Consort', 0.5235946178436279),
 ('queens', 0.5181134343147278),
 ('sultan', 0.5098592638969421),
 ('monarchy', 0.5087411999702454)]
```

Question 2.4: Finding Analogies [code + written] (2 Points)

Find an example of analogy that holds according to these vectors (i.e. the intended word is ranked top). In your solution please state the full analogy in the form x:y :: a:b. If you believe the analogy is complicated, explain why the analogy holds in one or two sentences.

Note: You may have to try many analogies to find one that works!

2.3 synonyms & antonyms 2 / 2

✓ - 0 pts Correct

- 1 pts Did not provide a possible reason
- 1 pts Did not provide valid triple

In [26]:

```
# -----
# Write your analogy exploration code here.

pprint.pprint(wv_from_bin.most_similar(positive=['man', 'actress'], negative=['woman']))
print()
pprint.pprint(wv_from_bin.most_similar(positive=['man', 'girl'], negative=['woman']))
print()
pprint.pprint(wv_from_bin.most_similar(positive=['large', 'low'], negative=['small']))

# -----
```

[('actor', 0.8202018141746521),
 ('Actor', 0.6390689611434937),
 ('funnyman', 0.6047691106796265),
 ('thespian', 0.5847228765487671),
 ('thesp', 0.5750559568405151),
 ('heart_throb', 0.5708590745925903),
 ('starlet', 0.5672707557678223),
 ('singer', 0.5657444000244141),
 ('Actress', 0.5615350008010864),
 ('ADRIEN_BRODY', 0.5556015372276306)]

[('boy', 0.8748372793197632),
 ('teenager', 0.6932151913642883),
 ('teenage_girl', 0.6277071237564087),
 ('lad', 0.624101996421814),
 ('kid', 0.6167631149291992),
 ('schoolboy', 0.600092351436615),
 ('teen_ager', 0.5852149724960327),
 ('boys', 0.5807873606681824),
 ('youngster', 0.5773991346359253),
 ('guy', 0.5459756255149841)]

[('high', 0.6629930734634399),
 ('Low', 0.5137864947319031),
 ('lower', 0.5012373924255371),
 ('higher', 0.49018242955207825),
 ('lowest', 0.48971161246299744),
 ('abysmally_low', 0.4538779556751251),
 ('highest', 0.4514465928077698),
 ('levels', 0.4280760586261749),
 ('rising', 0.4265591502189636),
 ('ahigh', 0.4251363277435303)]

Answer:

woman:actress::man:actor

woman:girl::man:boy

small:low::large:high

Question 2.5: Incorrect Analogy [code + written] (1 point)

Find an example of analogy that does *not* hold according to these vectors. In your solution, state the intended analogy in the form x:y :: a:b, and state the (incorrect) value of b according to the word vectors.

```
In [27]: # -----
# Write your incorrect analogy exploration code here.

pprint.pprint(wv_from_bin.most_similar(positive=['woman', 'doctor'], negative=['man']))

# -----
[('gynecologist', 0.7093890905380249),
 ('nurse', 0.6477286219596863),
 ('doctors', 0.6471461653709412),
 ('physician', 0.64389967918396),
 ('pediatrician', 0.6249487400054932),
 ('nurse_practitioner', 0.6218313574790955),
 ('obstetrician', 0.6072014570236206),
 ('ob_gyn', 0.5986713171005249),
 ('midwife', 0.5927063226699829),
 ('dermatologist', 0.5739566683769226)]
```

Answer:

Intended: man:doctor::woman:physician, actual: man:doctor::woman:gynecologist, gynecologist should be more generalized.

Question 2.6: Guided Analysis of Bias in Word Vectors [written] (1 point)

It's important to be cognizant of the biases (gender, race, sexual orientation etc.) implicit to our word embeddings.

Run the cell below, to examine (a) which terms are most similar to "woman" and "boss" and most dissimilar to "man", and (b) which terms are most similar to "man" and "boss" and most dissimilar to "woman". What do you find in the top 10?

2.4 correct analogy 2 / 2

- ✓ + 2 pts Valid analogy is top-1
- + 1 pts Valid analogy but not top-1
- + 1 pts Unclear analogy
- + 0 pts Incorrect/missing analogy

Answer:

woman:actress::man:actor

woman:girl::man:boy

small:low::large:high

Question 2.5: Incorrect Analogy [code + written] (1 point)

Find an example of analogy that does *not* hold according to these vectors. In your solution, state the intended analogy in the form x:y :: a:b, and state the (incorrect) value of b according to the word vectors.

```
In [27]: # -----
# Write your incorrect analogy exploration code here.

pprint.pprint(wv_from_bin.most_similar(positive=['woman', 'doctor'], negative=['man']))

# -----
[('gynecologist', 0.7093890905380249),
 ('nurse', 0.6477286219596863),
 ('doctors', 0.6471461653709412),
 ('physician', 0.64389967918396),
 ('pediatrician', 0.6249487400054932),
 ('nurse_practitioner', 0.6218313574790955),
 ('obstetrician', 0.6072014570236206),
 ('ob_gyn', 0.5986713171005249),
 ('midwife', 0.5927063226699829),
 ('dermatologist', 0.5739566683769226)]
```

Answer:

Intended: man:doctor::woman:physician, actual: man:doctor::woman:gynecologist, gynecologist should be more generalized.

Question 2.6: Guided Analysis of Bias in Word Vectors [written] (1 point)

It's important to be cognizant of the biases (gender, race, sexual orientation etc.) implicit to our word embeddings.

Run the cell below, to examine (a) which terms are most similar to "woman" and "boss" and most dissimilar to "man", and (b) which terms are most similar to "man" and "boss" and most dissimilar to "woman". What do you find in the top 10?

2.5 incorrect analogy 0 / 1

+ 1 pts Correct analogy, top-1 incorrect

✓ + 0 pts Incorrect analogy

+ 0 pts Correct analogy, top-1 correct

+ 0 pts Correct analogy, incorrect implementation/code

+ 0 pts Not attempted

```
In [30]: # Run this cell
# Here `positive` indicates the list of words to be similar to and `negative` indicates the list of words to be
# most dissimilar from.
pprint.pprint(wv_from_bin.most_similar(positive=['woman', 'boss'], ne
gative=['man']))
print()
pprint.pprint(wv_from_bin.most_similar(positive=['man', 'boss'], nega
tive=['woman']))
```

```
[('bosses', 0.552264392375946),
 ('manageress', 0.49151355028152466),
 ('exec', 0.45940810441970825),
 ('Manageress', 0.45598432421684265),
 ('receptionist', 0.4474116563796997),
 ('Jane_Danson', 0.44480547308921814),
 ('Fiz_Jennie_McAlpine', 0.44275766611099243),
 ('Coronation_Street_actress', 0.44275569915771484),
 ('supremo', 0.4409852623939514),
 ('coworker', 0.4398624897003174)]

[('supremo', 0.6097398400306702),
 ('MOTHERWELL_boss', 0.5489562153816223),
 ('CARETAKER_boss', 0.5375303626060486),
 ('Bully_Wee_boss', 0.533397376537323),
 ('YEOVIL_Town_boss', 0.5321705341339111),
 ('head_honcho', 0.5281979441642761),
 ('manager_Stan_Ternent', 0.5259714722633362),
 ('Viv_Busby', 0.5256163477897644),
 ('striker_Gabby_Agbonlahor', 0.5250812768936157),
 ('BARNESLEY_boss', 0.5238943696022034)]
```

Answer:

The prediction of positive=['woman', 'boss'], negative=['man']) is less appropriate than positive=['man', 'boss'], negative=['woman'], "receptionist", "Coronation_Street_actress", "coworker" are inappropriate even disrespectful like coronation_street_actress, bosses is complex of boss, however man and woman are all singular

Question 2.7: Independent Analysis of Bias in Word Vectors [code + written] (2 points)

Use the `most_similar` function to find another case where some bias is exhibited by the vectors. Please briefly explain the example of bias that you discover.

2.6 guided bias exploration 1 / 1

✓ - 0 pts Correct

- 1 pts Incorrect / Insufficient explanation

```
In [29]: # -----
# Write your bias exploration code here.

pprint.pprint(wv_from_bin.most_similar(positive=['woman', 'doctor'], negative=['man']))
print()
pprint.pprint(wv_from_bin.most_similar(positive=['man', 'doctor'], negative=['woman']))

# -----
```

[('gynecologist', 0.7093890905380249),
 ('nurse', 0.6477286219596863),
 ('doctors', 0.6471461653709412),
 ('physician', 0.64389967918396),
 ('pediatrician', 0.6249487400054932),
 ('nurse_practitioner', 0.6218313574790955),
 ('obstetrician', 0.6072014570236206),
 ('ob_gyn', 0.5986713171005249),
 ('midwife', 0.5927063226699829),
 ('dermatologist', 0.5739566683769226)]

[('physician', 0.646366536617279),
 ('doctors', 0.5858404636383057),
 ('surgeon', 0.572394073009491),
 ('dentist', 0.552364706993103),
 ('cardiologist', 0.5413815975189209),
 ('neurologist', 0.5271126627922058),
 ('neurosurgeon', 0.5249835848808289),
 ('urologist', 0.5247740149497986),
 ('Doctor', 0.5240625143051147),
 ('internist', 0.5183224081993103)]

Answer:

The prediction for positive=['man','doctor'], negative=['woman'] is much more accurate than positive=['woman','doctor'], negative=['man'], we expect man:doctor:woman:physician for both two above. Woman and man should have equal opportunity of being doctor/physician rather than just gynecologist at the most, which is not equal for woman and inappropriate, it's same problem as shown in question 2.6

Question 2.8: Thinking About Bias [written] (1 point)

What might be the cause of these biases in the word vectors?

2.7 independent bias exploration 2 / 2

✓ + 1 pts Found bias

✓ + 1 pts Explained bias

+ 0 pts Incorrect/Not attempted

+ 1 pts 2-way comparison not provided and hence cannot comment on bias.

Answer:

The cause of the biases may come from the corpus for training the word vector, some words have less context content existance consistent with the common human sense(correct answer) in the corpus, so it is similar like when training a machine leaning model with a small number of dataset may have bias. For example, in question Question 2.7, man:doctor::woman:gynecologist, however it should be physician as man:doctor::woman:physician, it may because woman have most context with gynecologist in the corpus may due to most of the OB doctor is woman. so the wv for gynecologist is closer than physician for woman

Submission Instructions

1. Click the Save button at the top of the Jupyter Notebook.
2. Please make a Gradescope account using your @stanford.edu email address (this is very important to help us enter your grade at the end of the quarter), and ensure your SUID is entered at the top of this notebook too.
3. Select Cell -> All Output -> Clear. This will clear all the outputs from all cells (but will keep the content of all cells).
4. Select Cell -> Run All. This will run all the cells in order, and will take several minutes.
5. Once you've rerun everything, select File -> Download as -> PDF via LaTeX
6. Look at the PDF file and make sure all your solutions are there, displayed correctly. The PDF is the only thing your graders will see!
7. Submit your PDF on Gradescope.

In []:

2.8 thinking about bias 1 / 1

✓ + 1 pts Correct: biases are picked up from the training corpus (Google News)

+ 0 pts Insufficient/missing explanation